

Final Project Progress Report (Beta)

Project Description

My project simulates a CPU scheduler. I plan to compare first come, first served (FCFS), shortest job first (SJF), shortest time remaining first (STRF), and round robin (RR) approaches by examining their CPU utilization, throughput, turnaround time, waiting time, and response time.

I use a loop with a counter to simulate the clock and have a queue of “processes,” structs that have metadata and an array of burst times to represent how much CPU time they get in one shot. The processes are rotated through, given “CPU time,” while collecting data about how long it takes, according to the loop clock.

Implementation

Interesting Code

simulator.c contains the simulation code. A set of “processes” is loaded and
schedules.c contains the comparison functions which embody the various scheduling approaches. Since the scheduling approach determines how tasks are prioritized and prioritization is manifested in how the processes are sorted in the heap, each comparison function corresponds to an approach.

Data Structures (I'm not expected to write these myself, right?)

heap.c Generic heap implementation for queuing jobs for CPU and IO.
Source: Jason Marcell (jasmarc) on GitHub

queue.c Generic queue for queuing processes before entering the heaps and after execution is completed. Implementation coming soon!

Other

expngen.c Exponential distribution generator to produce burst times. Will be modified to return single values and/or arrays rather than writing to a file.
Source: Ken Christensen at University of South Florida

Sample Output

The simulator will write meta data to files, which I will analyze and graph in a different program. The files will contain the information about CPU utilization (UT), throughput (THR), turnaround time (TRN), waiting time (WT), and response time (RSP). The simulator will run multiple (10-20) times for each approach, and each file will store the information for one approach, with each line of the file representing an execution of the simulator. Here's an example of an output file, with completely fake data:

UT	THR	TRN	WT	RSP
90	6.8	116	22	5
87	7.3	140	27	8
92	6.7	122	20	4
86	7.1	136	26	10
..

Current Progress

File	Done	To do
simulator	<ul style="list-style-type: none">- Process struct designed- Main loop partially implemented and otherwise pseudocoded/wireframed	<ul style="list-style-type: none">- Finish coding simulator- Writing to the output file
schedules	<ul style="list-style-type: none">- FCFS and SJF comparison functions written, but those were the easy ones	<ul style="list-style-type: none">- STRF and RR are preemptive approaches, so they will take more consideration- Make the function parameters void pointers, as the heap requires. Alternately, change the heap so that it is not generic
heap	<ul style="list-style-type: none">- Acquired	--
queue	--	<ul style="list-style-type: none">- Find an implementation
expgen	<ul style="list-style-type: none">- Acquired	<ul style="list-style-type: none">- Adapt to not write to file, but simply return a single value or array